

Hypothesis Testing

Cautions and Errors



Hypothesis Testing Steps

1. Assume the null hypothesis.
2. Find the probability of observing a sample at least as extreme as the sample you have if the null is true (**p-value**).
3. If the probability is low enough (below the **significance level**), reject the null in favor of the alternative.



Cautions about p -values

The use of p -values has become more controversial in recent years due to how often they are either misused or misunderstood.

See, for example, this Nature editorial:

<https://www.nature.com/articles/d41586-019-00874-8>

Cautions about p -values

Important:

- p -values do not give the likelihood that the result is due to chance
- p -values only summarize the data, assuming the null hypothesis is true! They do not say how likely the result is to be true.
- p -values say nothing about the size of an effect. Statistical significance is not the same as *practical* significance.
- A low p -value does not prove the alternative. Ronald Fisher, the inventor of the p -value, only meant for “statistical significance” to be an informal index.

Cautions about p -values

Another easy mistake to make with p -values is the **multiple comparisons/multiple testing** problem. When doing many simultaneous comparisons across a dataset, the chances increase of seeing a “statistically significant” effect which is just due to random sampling error.

See this xkcd comic: <https://xkcd.com/882/> or this FiveThirtyEight interactive:
<https://fivethirtyeight.com/features/science-isnt-broken/#part1>

Cautions about p -values

When doing hypothesis testing, it is important to distinguish between exploratory analysis and hypothesis testing.

Hypothesis testing must be deliberate, with a specific hypothesis in mind prior to looking at the data.

It is not valid to first look for potential effects in a dataset and then test those effects using the same data.

Hypothesis Testing - Types of Errors

		Reality	
		Null is True	Null is False
Our Decision	Do not Reject Null		
	Reject Null		



Hypothesis Testing - Types of Errors

		Reality	
		Null is True	Null is False
Our Decision	Do not Reject Null	Correct Decision	
	Reject Null		Correct Decision



Hypothesis Testing - Types of Errors

		Reality	
		Null is True	Null is False
Our Decision	Do not Reject Null	Correct Decision	
	Reject Null	False Positive / Type I Error	Correct Decision



NASHVILLE
SOFTWARE
SCHOOL

Hypothesis Testing - Types of Errors

How likely are we to make a Type I/False Positive error?

Note: A Type I error is only possible if the null is true.

Assuming that the null is true, how often will we commit a Type I error?

We reject the null hypothesis if the probability of a sample at least as extreme as ours is less than the significance level.

This happens with probability equal to the significance level (usually 5% of the time).

Hypothesis Testing - Types of Errors

		Reality	
		Null is True	Null is False
Our Decision	Do not Reject Null	Correct Decision	False Negative / Type II Error
	Reject Null	False Positive / Type I Error	Correct Decision



Hypothesis Testing - Types of Errors

How likely are we to make a Type II/False Negative error?

Note: A Type II error is only possible if the null is false.

Assuming that the null is false, how often will we commit a Type II error?

This depends on the effect size - how far off the true parameter is from the parameter value assumed by the null.

It also depends on how large a sample size we are using. At a larger sample size, we are more likely to detect a difference.

The probability of rejecting the null when it is false is the **power** of the test and can be influenced by carefully selecting a sample size.